

# “BetaGo”: Predicting Professional Players’ Moves in Go Game

Molly Zhang

---


# 1 minute guide to Go Game

## Three basic Rule:

- Black and white stone alternately put a stone on board (at intersection of two grid lines)
- If stone(s) are surrounded by opponent stones, it's removed from board
- Goal of game:  
whichever side that occupied more territory on board wins

Example: demo game

# Parsing Go game into Matrix



Go game board



```
1 (;SZ[19]FF[3]
2 PW[Sun Ce]
3 PB[Lu: Fan]
4 DT[196]
5 PC[Wu Kingdom, China]
6 RU[Old Chinese]
7 RE[? (No further moves known)]
8 GC[This is reputed to be the oldest game on
9 record. It is preserved in the Song Dynasty
10 go
11 manual "Carefree and Innocent Pastime" but
12 there are much earlier references to
13 famous Games of Wu.]
14 US[GoGoD95]
15 AB[pd][dp]
16 AW[dd][pp]
17 ;W[cn];B[qg];W[en];B[ci];W[ck];B[cf];W[ce];B[
18 df];W[gc];B[ic];W[iq];B[bo]
19 ;W[bn];B[lc];W[qf];B[qk];W[qi];B[qn];W[nc];B[
20 nd];W[qc];B[pc];W[qd];B[qb]
21 ;W[rb];B[pb];W[qo];B[pn];W[lm];B[cc];W[dc];B[
22 be];W[bd];B[cd];W[bc];B[de]
23 ;W[cb];B[rh];W[qg];B[ri];W[ho];B[eg];W[no]
24 )
```

Game recorded in "SGF" format



	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s
a	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
b	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	2	2	1
c	0	1	2	1	0	0	1	0	2	0	0	2	0	1	0	2	1	0	0
d	0	1	2	1	0	0	0	0	0	0	0	0	0	2	0	2	1	0	0
e	0	2	0	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
f	0	0	2	2	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
g	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
h	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0
i	0	0	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	2
j	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
k	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0
l	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
n	0	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	2	2	0
o	0	2	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0	0
p	0	0	0	2	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
q	0	0	0	2	0	2	0	1	0	0	0	1	0	0	0	0	0	0	0
r	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
s	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

Board as 19x19 matrix

n number of matrices from each game, n = total\_steps\_in\_game

# Data Summary

- Downloaded ~86,000 professional games in .sgf format from [gogodonline.co.uk](http://gogodonline.co.uk) (for \$15...)
- Each game as on average 200 moves, therefore giving a total of  $200 \times 86000 = 17$  million board positions and corresponding moves

## Question

- Given a board, with my stones and opponent stones on it, where should I go next?

# Overall Results:

- Best overall prediction accuracy over entire game is ~4%
- V.S. Accuracy of random guess (~0.42%)
- V.S. Accuracy of google AlphaGo (~50%)
  - What google DeepMind did and I didn't:

	BetaGo	AlphaGo
Training Size	1 million	30 million
Input layers (feature layer)	1	48
Conv Layers	1	many
Number of GPU and CPU	1 and 8	280 and 1920
Level of developer smartness	Okay	Very High


- However, my prediction accuracy of first 10-20 moves are high (~40%)

# Neural Net Models 1:


## One input layer, one hidden layer (ReLU activation) and one softmax output layer

**Input:**  
19 x 19 Board  
position matrix

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s
a	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
b	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	2	1	0
c	0	1	2	1	0	0	0	1	0	2	0	0	2	0	1	0	2	1	0
d	0	1	2	1	0	0	0	0	0	0	0	0	0	0	0	2	1	0	0
e	0	2	0	2	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0
f	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
g	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0
h	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0
i	0	0	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	2
j	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
k	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	2	0
l	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
m	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
n	0	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	2	2	0
o	0	2	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	1	0
p	0	0	0	0	2	0	0	0	0	0	0	0	0	0	0	0	1	0	0
q	0	0	0	0	0	2	0	2	0	1	0	0	0	0	0	0	0	0	0
r	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
s	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0



**Output:**  
Probability distribution  
of next move



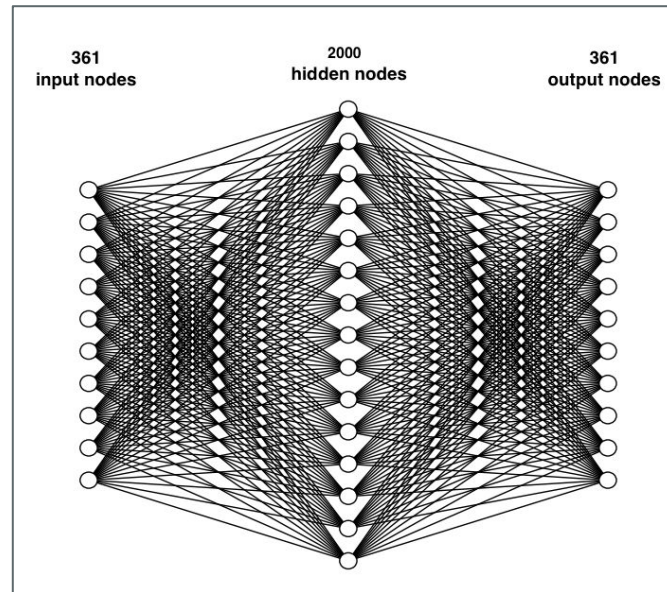


# Neural Net Models 1:

## Basic NN with one hidden layer and one softmax layer before output

### Experimented with:

- Different number of hidden nodes
- Accuracy decay with more stones on board
- Predicting accuracy scaling with increasing training data
- Regularization by dropping out
- Activation function: Sigmoid v.s. Rectifier



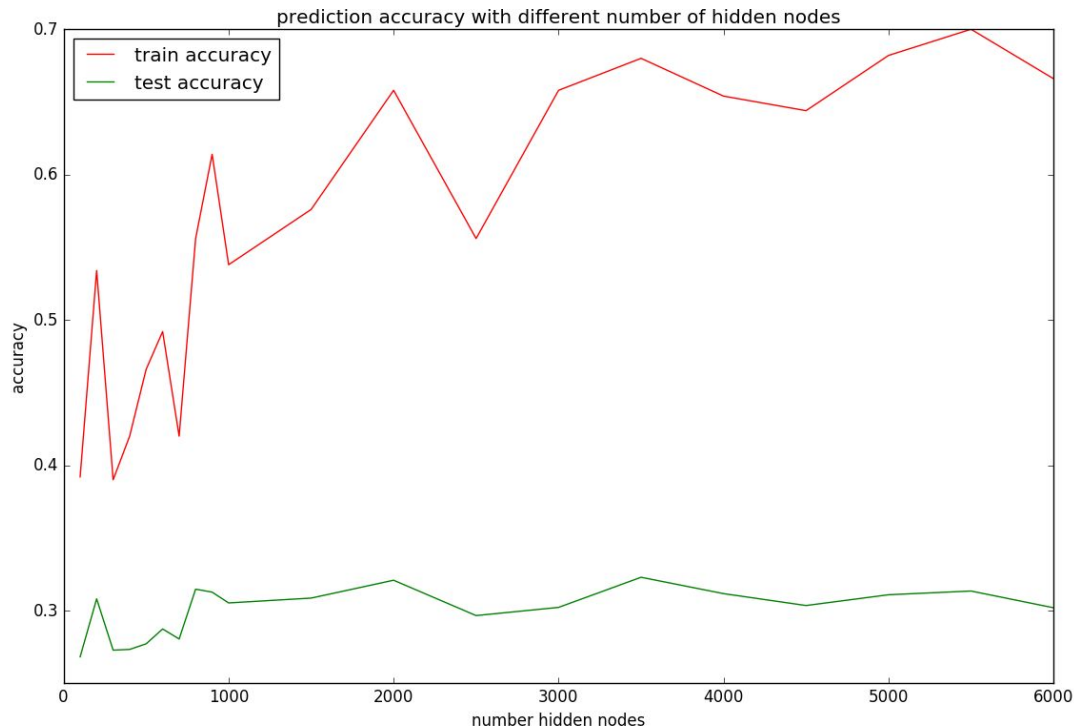
# Neural Net Models 1: hidden nodes on test accuracy and overfitting

## Take away:

- Lots of overfitting
- Best hidden nodes number: ~2000
- Tensorflow limits max node number to 6000

## Note:

Only 1000 training examples are used for speedy experiment (thus the lower accuracy and the high overfitting)

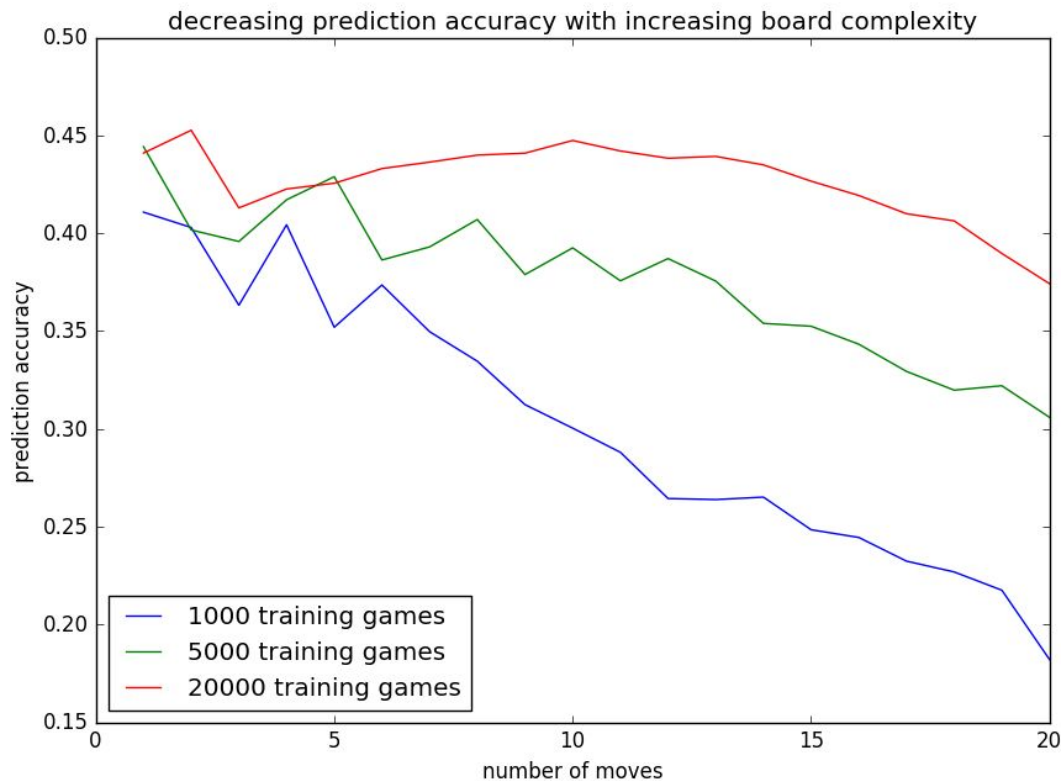




# Neural Net Models 1: Accuracy Decay with number stones on board

## Take away:

- As game progresses, it gets harder to predict
- More training data makes the accuracy drop slower

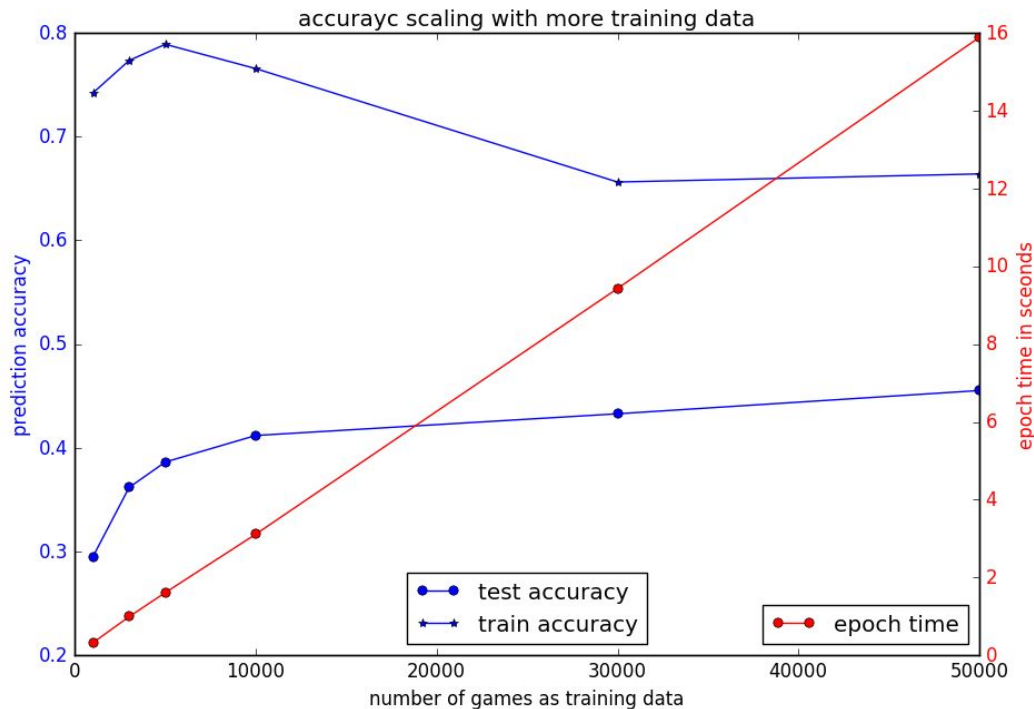


# Neural Net Models 1:

## Accuracy improvement and training slowdown with more training data

10,000 games is sweet trade-off between:

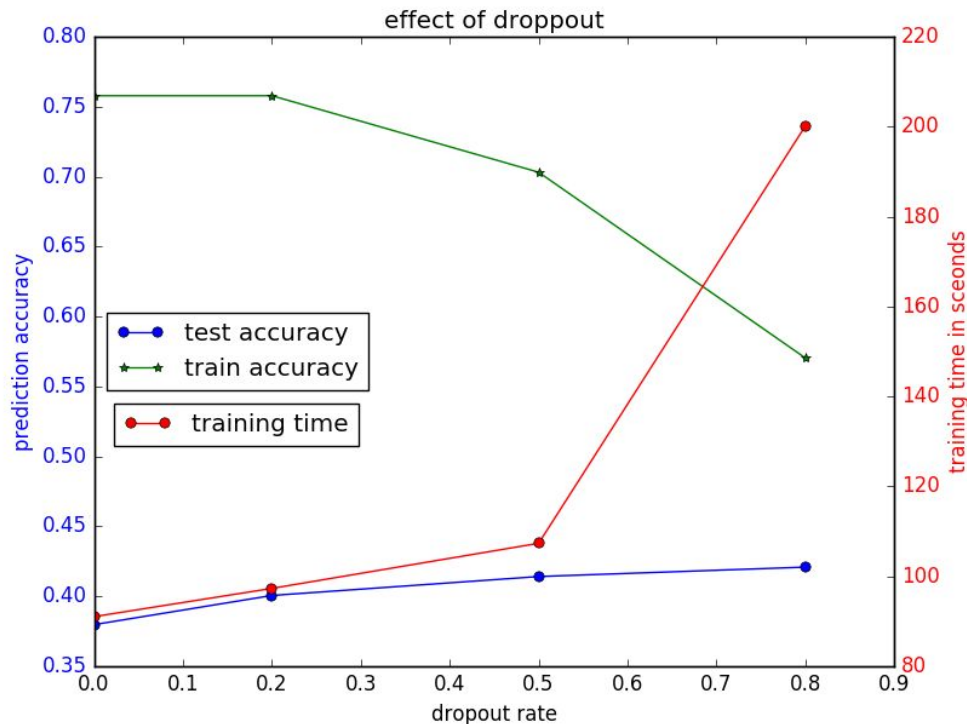
- Higher test accuracy
- Speedy epoch time




# Neural Net Models 1: Effect of regularization by dropping out

## Take away

- Dropout does reduce overfitting
- Dropout increase training time
- Dropout improves test accuracy
- Picking dropout rate = 0.5



## Neural Net Models 2: convolutional layer + fully connected layer + softmax output



Result:

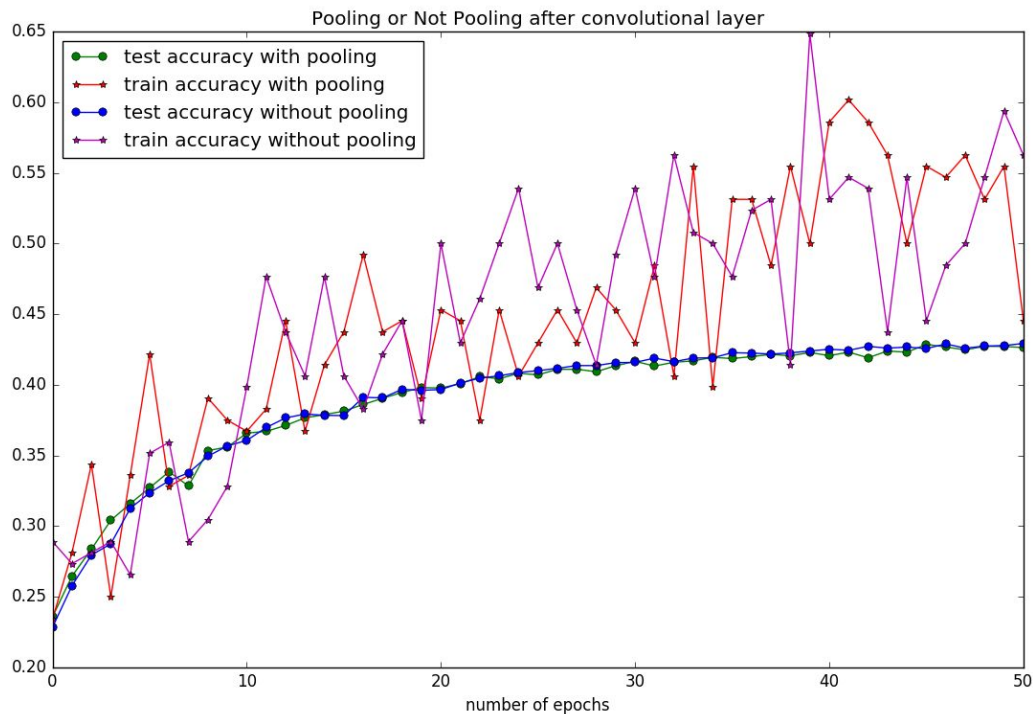
Accuracy: 43.3% (comparable to without convolutional layer)

Epoch time: 13 seconds (twice of NN model 1)

## Neural Net Models 2: Pooling, or No Pooling?

### Take away

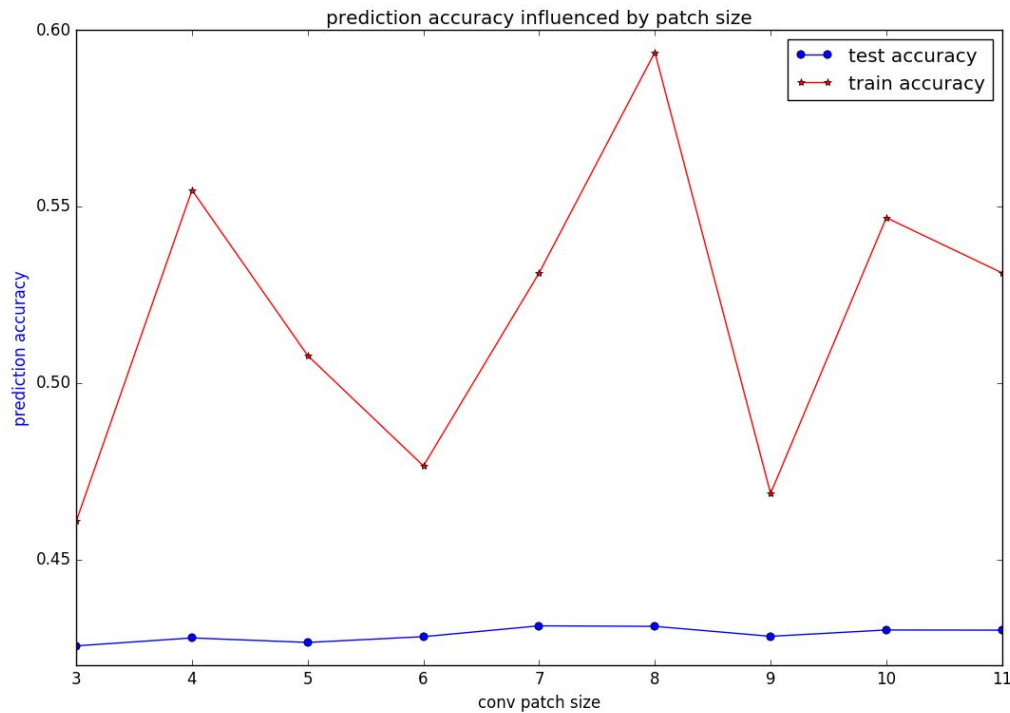
- Pooling reduces overfitting, by a little bit
- No difference in test accuracy
- Picking 2x2 max pooling because it reduces training time (and model complexity)





## Neural Net Models 2: What's the right convolutional patch size?

### Take away

- Patch size of 7 or 8 is best
- Matches gameplay intuition





# Visualize the Predictions in Game Play (o)







# Visualize the Predictions in Game Play (1)



# Visualize the Predictions in Game Play (3)



# Tools/Resources used

- Tensorflow (GPU version)
- Hyades (the supercomputing cluster in UCSC)

# References:

- Silver, David, et al. "Mastering the game of Go with deep neural networks and tree search." Nature 529.7587 (2016): 484-489.
- Clark, Christopher, and Amos Storkey. "Teaching deep convolutional neural networks to play go." arXiv preprint arXiv:1412.3409 (2014).

# Thanks

- David Parks for Inspiration, encouragement and discussions
- Keshav Mathur for GPU, cuda, batch loading etc help
- Dimitris Achlioptas for the class